

dplyr

The data frame is a key data structure in statistics and in R. The basic structure of a data frame is that there is one observation per row and each column represents a variable, a measure, feature, or characteristic of that observation. R has an internal implementation of data frames that is likely the one you will use most often. However, there are packages on CRAN that implement data frames via things like relational databases that allow you to operate on very very large data frames (but we won't discuss them here).

Given the importance of managing data frames, it's important that we have good tools for dealing with them.

Filtering, re-ordering, and collapsing, can often be tedious operations in R whose syntax is not very intuitive

The dplyr package is designed to mitigate a lot of these problems and to provide a highly optimized set of routines specifically for dealing with data frames.

The dplyr package was developed by Hadley Wickham of RStudio and is an optimized and distilled version of his plyr package. The dplyr package does not provide any “new” functionality to R per se, in the sense that everything dplyr does could already be done with base R, but it greatly simplifies existing functionality in R.

One important contribution of the dplyr package is that it provides a “grammar” (in particular, verbs) for data manipulation and for operating on data frames. With this grammar, you can sensibly communicate what it is that you are doing to a data frame that other people can understand (assuming they also know the grammar). This is useful because it provides an abstraction for data manipulation that previously did not exist. Another useful contribution is that the dplyr functions are very fast, as many key operations are coded in C++.

Grammar

Some of the key “verbs” provided by the dplyr package are

- **select**: return a subset of the columns of a data frame, using a flexible notation
- **filter**: extract a subset of rows from a data frame based on logical conditions
- **arrange**: reorder rows of a data frame
- **rename**: rename variables in a data frame
- **mutate**: add new variables/columns or transform existing variables
- **summarise** / **summarize**: generate summary statistics of different variables in the data frame, possibly within strata
- **%>%**: the “pipe” operator is used to connect multiple verb actions together into a pipeline

Common dplyr function properties:

1. The first argument is a data frame.
2. The subsequent arguments describe what to do with the data frame specified in the first argument, and you can refer to columns in the data frame directly without using the \$ operator (just use the column names).
3. The return result of a function is a new data frame

4. Data frames must be properly formatted and annotated for this to all be useful. In particular, the data must be tidy. In short, there should be one observation per row, and each column should represent a feature or characteristic of that observation.

```
library(dplyr)
```

```
##
## Attaching package: 'dplyr'
##
## The following objects are masked from 'package:stats':
##
##   filter, lag
##
## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union
```

```
select()
```

```
chicago <- readRDS("chicago.rds")
```

```
str(chicago)
```

```
## 'data.frame':   6940 obs. of  8 variables:
## $ city       : chr  "chic" "chic" "chic" "chic" ...
## $ tmpd       : num  31.5 33 33 29 32 40 34.5 29 26.5 32.5 ...
## $ dptp       : num  31.5 29.9 27.4 28.6 28.9 ...
## $ date       : Date, format: "1987-01-01" "1987-01-02" ...
## $ pm25tmean2: num   NA NA NA NA NA NA NA NA NA NA ...
## $ pm10tmean2: num   34 NA 34.2 47 NA ...
## $ o3tmean2   : num   4.25 3.3 3.33 4.38 4.75 ...
## $ no2tmean2  : num   20 23.2 23.8 30.4 30.3 ...
```

Suppose we wanted to take the first 3 columns only. There are a few ways to do this. We could for example use numerical indices. But we can also use the names directly.

```
names(chicago)[1:3]
```

```
## [1] "city" "tmpd" "dptp"
```

```
subset <- select(chicago, city:dptp)
head(subset)
```

```
##   city tmpd  dptp
## 1 chic 31.5 31.500
## 2 chic 33.0 29.875
## 3 chic 33.0 27.375
## 4 chic 29.0 28.625
## 5 chic 32.0 28.875
## 6 chic 40.0 35.125
```

we can omit variables using negating sign.

```
head(select(chicago, -(city:dptp)))
```

```
##           date pm25tmean2 pm10tmean2 o3tmean2 no2tmean2
## 1 1987-01-01           NA    34.00000 4.250000  19.98810
## 2 1987-01-02           NA           NA 3.304348  23.19099
## 3 1987-01-03           NA    34.16667 3.333333  23.81548
```

```
## 4 1987-01-04      NA    47.00000 4.375000 30.43452
## 5 1987-01-05      NA      NA 4.750000 30.33333
## 6 1987-01-06      NA    48.00000 5.833333 25.77233
```

specification of variablenames using patterns

```
subset <- select(chicago, ends_with("2"))
str(subset)
```

```
## 'data.frame':    6940 obs. of  4 variables:
## $ pm25tmean2: num  NA NA NA NA NA NA NA NA NA NA ...
## $ pm10tmean2: num  34 NA 34.2 47 NA ...
## $ o3tmean2 : num  4.25 3.3 3.33 4.38 4.75 ...
## $ no2tmean2 : num  20 23.2 23.8 30.4 30.3 ...
```

```
subset <- select(chicago, starts_with("d"))
str(subset)
```

```
## 'data.frame':    6940 obs. of  2 variables:
## $ dptp: num  31.5 29.9 27.4 28.6 28.9 ...
## $ date: Date, format: "1987-01-01" "1987-01-02" ...
```

filter()

The filter() function is used to extract subsets of rows from a data frame. This function is similar to the existing subset() function in R but is quite a bit faster.

Suppose we wanted to extract the rows of the chicago data frame where the levels of PM2.5 are greater than 30 (which is a reasonably high level), we could do

```
chic.f <- filter(chicago, pm25tmean2 > 30)
str(chic.f)
```

```
## 'data.frame':    194 obs. of  8 variables:
## $ city      : chr  "chic" "chic" "chic" "chic" ...
## $ tmpd      : num  23 28 55 59 57 57 75 61 73 78 ...
## $ dptp      : num  21.9 25.8 51.3 53.7 52 56 65.8 59 60.3 67.1 ...
## $ date      : Date, format: "1998-01-17" "1998-01-23" ...
## $ pm25tmean2: num  38.1 34 39.4 35.4 33.3 ...
## $ pm10tmean2: num  32.5 38.7 34 28.5 35 ...
## $ o3tmean2 : num  3.18 1.75 10.79 14.3 20.66 ...
## $ no2tmean2 : num  25.3 29.4 25.3 31.4 26.8 ...
```

```
summary(chic.f$pm25tmean2)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##  30.05   32.12   35.04   36.63   39.53   61.50
```

We can place an arbitrarily complex logical sequence inside of filter(), so we could for example extract the rows where PM2.5 is greater than 30 and temperature is greater than 80 degrees Fahrenheit.

```
chic.f <- filter(chicago, pm25tmean2 > 30 & tmpd > 80)
select(chic.f, date, tmpd, pm25tmean2)
```

```
##      date tmpd pm25tmean2
## 1 1998-08-23   81  39.60000
## 2 1998-09-06   81  31.50000
## 3 2001-07-20   82  32.30000
## 4 2001-08-01   84  43.70000
## 5 2001-08-08   85  38.83750
```

```
## 6 2001-08-09 84 38.20000
## 7 2002-06-20 82 33.00000
## 8 2002-06-23 82 42.50000
## 9 2002-07-08 81 33.10000
## 10 2002-07-18 82 38.85000
## 11 2003-06-25 82 33.90000
## 12 2003-07-04 84 32.90000
## 13 2005-06-24 86 31.85714
## 14 2005-06-27 82 51.53750
## 15 2005-06-28 85 31.20000
## 16 2005-07-17 84 32.70000
## 17 2005-08-03 84 37.90000
```

arrange()

The `arrange()` function is used to reorder rows of a data frame according to one of the variables/columns. Reordering rows of a data frame (while preserving corresponding order of other columns) is normally a pain to do in R. The `arrange()` function simplifies the process quite a bit.

Here we can order the rows of the data frame by date

```
chicago <- arrange(chicago, date)
head(select(chicago, date, pm25tmean2), 3)
```

```
##           date pm25tmean2
## 1 1987-01-01          NA
## 2 1987-01-02          NA
## 3 1987-01-03          NA
```

```
tail(select(chicago, date, pm25tmean2), 3)
```

```
##           date pm25tmean2
## 6938 2005-12-29    7.45000
## 6939 2005-12-30   15.05714
## 6940 2005-12-31   15.00000
```

Columns can be arranged in descending order too by using the special `desc()` operator

```
chicago <- arrange(chicago, desc(date))
head(select(chicago, date, pm25tmean2), 3)
```

```
##           date pm25tmean2
## 1 2005-12-31   15.00000
## 2 2005-12-30   15.05714
## 3 2005-12-29    7.45000
```

rename()

Renaming a variable in a data frame in R is surprisingly hard to do! The `rename()` function is designed to make this process easier.

```
head(chicago[, 1:5], 3)
```

```
##   city tmpd dptp      date pm25tmean2
## 1 chic  35 30.1 2005-12-31   15.00000
## 2 chic  36 31.0 2005-12-30   15.05714
## 3 chic  35 29.4 2005-12-29    7.45000
```

The `dptp` column is supposed to represent the dew point temperature and the `pm25tmean2` column provides the PM2.5 data. However, these names are pretty obscure or awkward and probably be renamed to something more sensible.

```
chicago <- rename(chicago, dewpoint = dptp, pm25 = pm25tmean2)
head(chicago[, 1:5], 3)
```

```
##   city tmpd dewpoint      date    pm25
## 1 chic   35    30.1 2005-12-31 15.00000
## 2 chic   36    31.0 2005-12-30 15.05714
## 3 chic   35    29.4 2005-12-29  7.45000
```

The syntax inside the `rename()` function is to have the new name on the left-hand side of the `=` sign and the old name on the right-hand side.

mutate()

The `mutate()` function exists to compute transformations of variables in a data frame. Often, you want to create new variables that are derived from existing variables and `mutate()` provides a clean interface for doing that.

For example, with air pollution data, we often want to detrend the data by subtracting the mean from the data. That way we can look at whether a given day's air pollution level is higher than or less than average (as opposed to looking at its absolute level).

```
chicago <- mutate(chicago, pm25detrend = pm25 - mean(pm25, na.rm = TRUE))
head(chicago)
```

```
##   city tmpd dewpoint      date    pm25 pm10tmean2  o3tmean2 no2tmean2
## 1 chic   35    30.1 2005-12-31 15.00000      23.5  2.531250  13.25000
## 2 chic   36    31.0 2005-12-30 15.05714      19.2  3.034420  22.80556
## 3 chic   35    29.4 2005-12-29  7.45000      23.5  6.794837  19.97222
## 4 chic   37    34.5 2005-12-28 17.75000      27.5  3.260417  19.28563
## 5 chic   40    33.6 2005-12-27 23.56000      27.0  4.468750  23.50000
## 6 chic   35    29.6 2005-12-26  8.40000       8.5 14.041667  16.81944
##   pm25detrend
## 1   -1.230958
## 2   -1.173815
## 3   -8.780958
## 4    1.519042
## 5    7.329042
## 6   -7.830958
```

There is also the related `transmute()` function, which does the same thing as `mutate()` but then drops all non-transformed variables.

Here we detrend the PM10 and ozone (O3) variables.

```
head(transmute(chicago,
               pm10detrend = pm10tmean2 - mean(pm10tmean2, na.rm = TRUE),
               o3detrend = o3tmean2 - mean(o3tmean2, na.rm = TRUE)))
```

```
##   pm10detrend  o3detrend
## 1 -10.395206 -16.904263
## 2 -14.695206 -16.401093
## 3 -10.395206 -12.640676
## 4  -6.395206 -16.175096
## 5  -6.895206 -14.966763
## 6 -25.395206  -5.393846
```

group-by() and summarize()

The `group_by()` function is used to generate summary statistics from the data frame within strata defined by a variable.

For example, in this air pollution dataset, you might want to know what the average annual level of PM2.5 is. So the stratum is the year, and that is something we can derive from the date variable. In conjunction with the `group_by()` function we often use the `summarize()` function (or `summarise()` for some parts of the world).

The general operation here is a combination of splitting a data frame into separate pieces defined by a variable or group of variables **`group_by()`**, and then applying a summary function across those subsets **`summarize()`**.

First, we can create a year variable using `as.POSIXlt()`

```
?as.POSIXlt
chicago <- mutate(chicago, year = as.POSIXlt(date)$year + 1900)
```

Now we can create a separate data frame that splits the original data frame by year.

```
years <- group_by(chicago, year)
```

Finally, we compute summary statistics for each year in the data frame with the `summarize()` function.

```
summarize(years, pm25 = mean(pm25, na.rm = TRUE),
           o3 = max(o3tmean2, na.rm = TRUE),
           no2 = median(no2tmean2, na.rm = TRUE))
```

```
## # A tibble: 19 x 4
##   year pm25    o3    no2
##   <dbl> <dbl> <dbl> <dbl>
## 1 1987 NaN    63.0  23.5
## 2 1988 NaN    61.7  24.5
## 3 1989 NaN    59.7  26.1
## 4 1990 NaN    52.2  22.6
## 5 1991 NaN    63.1  21.4
## 6 1992 NaN    50.8  24.8
## 7 1993 NaN    44.3  25.8
## 8 1994 NaN    52.2  28.5
## 9 1995 NaN    66.6  27.3
## 10 1996 NaN    58.4  26.4
## 11 1997 NaN    56.5  25.5
## 12 1998 18.3   50.7  24.6
## 13 1999 18.5   57.5  24.7
## 14 2000 16.9   55.8  23.5
## 15 2001 16.9   51.8  25.1
## 16 2002 15.3   54.9  22.7
## 17 2003 15.2   56.2  24.6
## 18 2004 14.6   44.5  23.4
## 19 2005 16.2   58.8  22.6
```

`summarize()` returns a data frame with year as the first column, and then the annual averages of pm25, o3, and no2.

In a slightly more complicated example, we might want to know what are the average levels of ozone (o3) and nitrogen dioxide (no2) within quintiles of pm25. A slicker way to do this would be through a regression model, but we can actually do this quickly with `group_by()` and `summarize()`

First, we can create a categorical variable of pm25 divided into quintiles.

```
qq <- quantile(chicago$pm25, seq(0, 1, 0.2), na.rm = TRUE)
qq
```

```
##      0%      20%      40%      60%      80%     100%
##  1.700  8.700 12.375 16.700 22.610 61.500
```

```
chicago <- mutate(chicago, pm25.quint = cut(pm25, qq))
```

Now we can group the data frame by the pm25.quint variable.

```
quint <- group_by(chicago, pm25.quint)
```

Finally, we can compute the mean of o3 and no2 within quintiles of pm25.

```
summarize(quint, o3 = mean(o3tmean2, na.rm = TRUE),
           no2 = mean(no2tmean2, na.rm = TRUE))
```

```
## # A tibble: 6 x 3
##   pm25.quint    o3    no2
##   <fct>      <dbl> <dbl>
## 1 (1.7,8.7]    21.7  18.0
## 2 (8.7,12.4]   20.4  22.1
## 3 (12.4,16.7]  20.7  24.4
## 4 (16.7,22.6]  19.9  27.3
## 5 (22.6,61.5]  20.3  29.6
## 6 <NA>        18.8  25.8
```

finally and again pipes

The pipeline operator %>% is very handy for stringing together multiple dplyr functions in a sequence of operations. Notice above that every time we wanted to apply more than one function, the sequence gets buried in a sequence of nested function calls that is difficult to read, i.e. > third(second(first(x)))

This nesting is not a natural way to think about a sequence of operations. The %>% operator allows you to string operations in a left-to-right fashion, i.e.

```
chicago %>%
  mutate(pm25.quint = cut(pm25, qq)) %>%
  group_by(pm25.quint) %>%
  summarize(o3 = mean(o3tmean2, na.rm = TRUE),
           no2 = mean(no2tmean2, na.rm = TRUE))
```

```
## # A tibble: 6 x 3
##   pm25.quint    o3    no2
##   <fct>      <dbl> <dbl>
## 1 (1.7,8.7]    21.7  18.0
## 2 (8.7,12.4]   20.4  22.1
## 3 (12.4,16.7]  20.7  24.4
## 4 (16.7,22.6]  19.9  27.3
## 5 (22.6,61.5]  20.3  29.6
## 6 <NA>        18.8  25.8
```

Another example might be computing the average pollutant level by month. This could be useful to see if there are any seasonal trends in the data.

```
chicago %>%
  mutate(chicago, month = as.POSIXlt(date)$mon + 1) %>%
  group_by(month) %>%
```

```
summarize(pm25 = mean(pm25, na.rm = TRUE),
          o3 = max(o3tmean2, na.rm = TRUE),
          no2 = median(no2tmean2, na.rm = TRUE))
```

```
## # A tibble: 12 x 4
##   month pm25    o3    no2
##   <dbl> <dbl> <dbl> <dbl>
## 1     1    17.8  28.2  25.4
## 2     2    20.4  37.4  26.8
## 3     3    17.4  39.0  26.8
## 4     4    13.9  47.9  25.0
## 5     5    14.1  52.8  24.2
## 6     6    15.9  66.6  25.0
## 7     7    16.6  59.5  22.4
## 8     8    16.9  54.0  23.0
## 9     9    15.9  57.5  24.5
## 10    10    14.2  47.1  24.2
## 11    11    15.2  29.5  23.6
## 12    12    17.5  27.7  24.5
```

Summary

The dplyr package provides a concise set of operations for managing data frames. With these functions we can do a number of complex operations in just a few lines of code. In particular, we can often conduct the beginnings of an exploratory analysis with the powerful combination of `group_by()` and `summarize()`

Once you learn the dplyr grammar there are a few additional benefits

1. dplyr can work with other data frame “backends” such as SQL databases. There is an SQL interface for relational databases via the DBI package
2. dplyr can be integrated with the data.table package for large fast tables

The dplyr package is handy way to both simplify and speed up your data frame management code. It’s rare that you get such a combination at the same time!